Maya Fernandes

PhD Candidate, Machine Learning

Edinburgh, United Kingdom | maya.fernandes@example.ac.uk
github.com/mayaf-example | scholar.example.com/mayaf | mayafernandes.example.com

RESEARCH INTERESTS

Robustness of machine learning under distribution shift; efficient inference and evaluation for large language models; open, reproducible benchmarks.

EDUCATION

Westbrook University Sep 2022 – Expected Jun 2027
Ph.D. in Computer Science *Edinburgh, UK*
Advisor: Prof. A. Lindqvist · Thesis: Measuring and mitigating distribution shift in deployed language models

Westbrook University Sep 2021 – Aug 2022
M.Sc. in Artificial Intelligence (Distinction) *Edinburgh, UK*

University of Lisbourne Sep 2017 – Jun 2021
B.Sc. in Mathematics and Computer Science *Lisbon, Portugal*
First-class honours

PUBLICATIONS

1. **M. Fernandes**, R. Okafor, A. Lindqvist. "ShiftBench: A Unified Benchmark for Distribution Shift in Language Model Evaluation." *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2025. arxiv.org/abs/0000.00001
2. J. Chen, **M. Fernandes**, A. Lindqvist. "Calibrated Early Exiting for Efficient Transformer Inference." *International Conference on Learning Representations (ICLR)*, 2025.
3. **M. Fernandes**, A. Lindqvist. "When Does Fine-Tuning Hurt Robustness? An Empirical Study." *Workshop on Distribution Shifts, NeurIPS*, 2023. Oral presentation.

RESEARCH & PROFESSIONAL EXPERIENCE

Research Scientist Intern Jun 2024 – Sep 2024
Northlight AI Research *London, UK*

- Developed a shift-detection method for production LLM traffic that flagged 92% of injected drift events with a 3% false-positive rate.
- Released the evaluation harness internally; adopted by 2 product teams for weekly model monitoring.

RESEARCH SOFTWARE

shiftbench | github.com/mayaf-example/shiftbench Jan 2024 – Present
Python, PyTorch, Hugging Face

- Open-source benchmark library accompanying the NeurIPS paper: 14 datasets, 30+ baselines, 350 GitHub stars, 9 external contributors.

early-exit-transformers | github.com/mayaf-example/early-exit-transformers Sep 2024 – Present
PyTorch, CUDA

- Reference implementation for the ICLR paper; 1.8x inference speed-up at equal accuracy on GLUE.

HONORS & AWARDS

Doctoral Training Studentship (full funding) — Westbrook University Sep 2022

Best Paper Award — Workshop on Distribution Shifts, NeurIPS Dec 2023

TEACHING

Teaching Assistant, Machine Learning Practical 2023 – 2025
Westbrook University · 180 students

Guest Lecture, Robustness in Deep Learning Mar 2025
Westbrook University, MSc AI

INVITED TALKS

Benchmarks that age well Feb 2026
ML Reading Group, University of Lisbourne

Measuring shift in deployed LLMs Sep 2024
Northlight AI Research

SERVICE

- Reviewer: NeurIPS (2024, 2025), ICLR (2025, 2026), TMLR
- Organiser, Westbrook ML Reading Group (2023–present)

TECHNICAL SKILLS

Programming: Python, PyTorch, JAX, CUDA, C++, SQL

ML tooling: Hugging Face, Weights & Biases, Slurm, Docker

LANGUAGES

English (Fluent) | Portuguese (Native)